


# Semaphores

CS511

# Motivation

- ▶ Algorithms for mutex seen up until now run on any machine (they only use standard instructions)
- ▶ These are too low-level to be used reliably
- ▶ Semaphores are higher-level constructs
  - ▶ Usually implemented by the OS
  - ▶ Widely used in many PLs

# States of a Process



- ▶ A scheduler decides which of the ready processes it should run
  - ▶ Arbitrary interleavings = we assume nothing about the scheduler

# Semaphore

A semaphore is an Abstract Data Type with:

- ▶ (Atomic!) Operations:
  - ▶ acquire (or wait)
  - ▶ release (or signal)
- ▶ Data fields:
  - ▶ permissions: non-negative integer
  - ▶ processes: set of processes

# Acquire

Acquire consumes a permission, waits if none are available

```
1  atomic acquire() {  
2      currentThread = Thread.currentThread();  
3      if (permissions > 0) {  
4          permissions--;  
5      } else {  
6          processes.add(thread);  
7          currentThread.state = BLOCKED;  
8      }  
9  }
```

# Release

Release frees a permission (wakens a blocked thread, if there are any)

```
1  atomic release() {
2      if (processes.empty()) {
3          permissions++;
4      } else {
5          wakingThread = processes.removeAny();
6          wakingThread.state = READY;
7      }
8  }
```

# Mutex or Binary Semaphore

- ▶ A semaphore that only admits 0 or 1 permissions.
  - ▶ Semaphores that allow arbitrary values of permission are called **counting** semaphores
- ▶ Initialized to  $(0, \emptyset)$  or  $(1, \emptyset)$
- ▶ The acquire operation is unchanged
- ▶ The release operation is now defined as:

```
1  atomic release() {
2      if (permissions == 1) {
3          // do nothing
4      } else if (processes.empty()) {
5          permissions = 1;
6      } else {
7          wakingThread = processes.removeAny();
8          wakingThread.state = READY;
9      }
10 }
```

Note: if permissions is 1, successive calls to `release` are lost

# Mutual Exclusion using mutex

- ▶ We now revisit the critical section problem for two processes
- ▶ Using a semaphore `mutex` the solution is trivial
  - ▶ P wishes to enter: preprotocol `mutex.acquire()`
  - ▶ When P exits: postprotocol `mutex.release()`



# Mutual Exclusion using mutex

```
1 global Semaphore mutex = new Semaphore(1);

1 thread P: {
2     // non-critical section
3     mutex.acquire();
4     // critical section
5     mutex.release();
6     // non-critical section
7 }

1 thread Q: {
2     // non-critical section
3     mutex.acquire();
4     // critical section
5     mutex.release();
6     // non-critical section
7 }
```

# Mutual Exclusion using mutex

```
1 global Semaphore mutex = new Semaphore(1);

1 thread P: {                                1 thread Q: {
2     // non-critical section                2     // non-critical section
3     mutex.acquire();                       3     mutex.acquire();
4     // critical section                    4     // critical section
5     mutex.release();                       5     mutex.release();
6     // non-critical section                6     // non-critical section
7 }                                           7 }
```

This solution does not use busy waiting since a process that blocks in the acquire goes into the BLOCKED state and only returns to the READY state once it is given permission to do so.

# Semaphores in Java

Class Semaphore in `java.util.concurrent`

► `java.util.concurrent.Semaphore`

```
1  /** Creates a semaphore with the given number of permits */
2  Semaphore(int permits)

1  /** Acquires a permit from this Semaphore,
2      blocking until one is available */
3  void acquire()

1  /** Releases a permit, returning it to the semaphore */
2  void release()
```

# Semaphores in Java

Class Semaphore in java.util.concurrent

► java.util.concurrent.Semaphore

```
1  /** Creates a semaphore with the given number of permits */
2  Semaphore(int permits)

1  /** Acquires a permit from this Semaphore,
2      blocking until one is available */
3  void acquire()

1  /** Releases a permit, returning it to the semaphore */
2  void release()
```

Example:

```
1  Semaphore mutex = new Semaphore(1);
2  mutex.acquire()
3  // critical section
4  mutex.release()
```

# Semaphore Invariants

Let  $k$  be the initial value of the `permissions` field of a semaphore  $s$

where

- ▶ `#releases` is the number of `s.release()` statements executed
- ▶ `#acquire` is the number of `s.acquire()` statements executed
- ▶ A blocked process is considered not to have executed an `acquire` operation.

# Semaphore Invariants

Let  $k$  be the initial value of the `permissions` field of a semaphore  $s$

1.  $\text{permissions} \geq 0$

where

- ▶ `#releases` is the number of `s.release()` statements executed
- ▶ `#acquire` is the number of `s.acquire()` statements executed
- ▶ A blocked process is considered not to have executed an `acquire` operation.

# Semaphore Invariants

Let  $k$  be the initial value of the `permissions` field of a semaphore  $s$

1.  $\text{permissions} \geq 0$
2.  $\text{permissions} = k + \text{\#releases} - \text{\#acquires}$

where

- ▶ `\#releases` is the number of `s.release()` statements executed
- ▶ `\#acquire` is the number of `s.acquire()` statements executed
- ▶ A blocked process is considered not to have executed an `acquire` operation.

# The Semaphore Solution for the MEP ( $k = 1$ )

`#criticalSection`: number of processes in their critical sections



# The Semaphore Solution for the MEP ( $k = 1$ )

`#criticalSection`: number of processes in their critical sections

1. `#criticalSection + permissions = 1`

# The Semaphore Solution for the MEP ( $k = 1$ )

`#criticalSection`: number of processes in their critical sections

1. `#criticalSection + permissions = 1`
2. `#criticalSection = #acquires - #releases`

# The Semaphore Solution for the MEP ( $k = 1$ )

`#criticalSection`: number of processes in their critical sections

1. `#criticalSection + permissions = 1`
2. `#criticalSection = #acquires - #releases`

Item 1 guarantees:

# The Semaphore Solution for the MEP ( $k = 1$ )

`#criticalSection`: number of processes in their critical sections

1. `#criticalSection + permissions = 1`
2. `#criticalSection = #acquires - #releases`

Item 1 guarantees:

- ▶ **Mutual Exclusion** (`#criticalSection  $\leq$  1` since `0  $\leq$  permissions`)

# The Semaphore Solution for the MEP ( $k = 1$ )

`#criticalSection`: number of processes in their critical sections

1. `#criticalSection + permissions = 1`
2. `#criticalSection = #acquires - #releases`

Item 1 guarantees:

- ▶ **Mutual Exclusion** (`#criticalSection`  $\leq 1$  since  $0 \leq \text{permissions}$ )
- ▶ **Absence of deadlock** (it never happens that `permissions = 0` and `#criticalSection = 0`)

# The Semaphore Solution for the MEP ( $k = 1$ )

`#criticalSection`: number of processes in their critical sections

1. `#criticalSection + permissions = 1`
2. `#criticalSection = #acquires - #releases`

Item 1 guarantees:

- ▶ **Mutual Exclusion** (`#criticalSection`  $\leq 1$  since  $0 \leq \text{permissions}$ )
- ▶ **Absence of deadlock** (it never happens that `permissions = 0` and `#criticalSection = 0`)
- ▶ **No starvation** between two processes

# The Turnstile Problem using Binary Semaphores

```
1 global int counter = 0;
2 global Semaphore mutex = new Semaphore(1);
3
4 turnstile() {
5     repeat (50) {
6         mutex.acquire();
7         counter++;
8         mutex.release();
9     }
10 }
11
12 repeat (2)
13     thread turnstile();
```

# The Turnstile Problem using Binary Semaphores (Java)

```
1 public class Turnstile extends Thread {
2     static volatile int counter = 0;
3     static Semaphore mutex = new Semaphore(1);
4     public void run() {
5         for(int i = 0; i < 50; i++){
6             mutex.acquire();
7             counter++;
8             mutex.release();
9             System.out.println(id+"- In comes: "+i );
10        }
11    }
12
13    public static void main(String args[]) {
14        try{
15            Thread m1 = new Turnstile(1);
16            m1.start();
17            Thread m2 = new Turnstile(2);
18            m2.start();
19        } catch(Exception e){}
20    }
21 }
```



# Counting Example in Java using Semaphores

```
1 public class Turnstile extends Thread {  
2     static volatile int counter = 0;  
3     ...  
}
```

- ▶ The `volatile` keyword is recommended for variables that are shared
- ▶ It guarantees that
  - ▶ Its value will never be cached thread-locally: all reads and writes will go straight to “main memory”; and
  - ▶ Access to the variable acts as though it is enclosed in a synchronized block, synchronized on itself (more later).

# Strong Semaphores

The same solution above for the critical section also works for  $N$  processes

```
1 repeat (N) {  
2     thread {  
3         // non-critical section  
4         mutex.acquire();  
5         // critical section  
6         mutex.release();  
7         // non-critical section  
8     }  
9 }
```

- ▶ But there is the possibility of **starvation**.
- ▶ The problem is caused by the fact that blocked processes are placed in a **set** of processes

# Strong Semaphores

- ▶ This can be remedied by changing the set to be a queue
- ▶ In Java this is indicated by the second argument of the constructor

```
1  /** Creates a Semaphore with the given number of permits
2      and the given fairness setting. */
3  Semaphore(int permits, boolean fair)
```

- ▶ When fairness is set to `true`, the semaphore gives permits to access mutual resources in the order the threads have asked for it (FIFO)

Semaphores

Synchronization Among Processes

# Synchronization Problems

- ▶ The **critical section problem** is an abstraction of the synchronization problems that occur when multiple processes compete for the same resource
- ▶ Another type of synchronization problem is when processes must coordinate the **order of execution**

# Revisiting the Turnstile Problem

Suppose we wish to print the counter total for  $N$  turnstiles

```
1 repeat (N)
2   thread turnstile();
3   print("Total = " + counter);
```

# Revisiting the Turnstile Problem

Suppose we wish to print the counter total for  $N$  turnstiles

```
1 repeat (N)
2   thread turnstile();
3   print("Total = " + counter);
```

What happens when we run this code?

# Revisiting the Turnstile Problem

```
1 global int counter = 0;
2 global Semaphore mutex = new Semaphore(1);
3 global Semaphore finish = new Semaphore(-N+1);
4
5 turnstile() {
6     repeat (100) {
7         mutex.acquire();
8         counter++;
9         mutex.release();
10    }
11    finish.release();
12 }
13
14 repeat (N)
15     thread turnstile();
16     finish.acquire();
17     print("Total = " + counter);
```



# Dining Philosophers



# Dining Philosophers



- Philosophers think and eat, in turns

# Dining Philosophers



- ▶ Philosophers think and eat, in turns
- ▶ They can only eat if they have both forks

# Dining Philosophers



- ▶ Philosophers think and eat, in turns
- ▶ They can only eat if they have both forks
- ▶ They can only grab the forks to their left and right

# Dining Philosophers

```
1  Philosopher(id) {  
2      while (true)  
3          // think  
4          // pick forks  
5          // eat  
6          // leave forks  
7  }
```

# Dining Philosophers

```
1  Philosopher(id) {  
2      while (true)  
3          // think  
4          // pick forks  
5          // eat  
6          // leave forks  
7  }
```

- ▶ Shared resource: the forks
- ▶ Mutex: at any given moment only one philosopher can have a fork
- ▶ Synchronization: a philosopher can only eat if she/he has both forks
- ▶ Absence of deadlock, livelock and starvation

## Dining Philosophers (naive attempt)

```
1 global Semaphore[] forks = [1,...,1]; // N
2
3 Philosopher(id) {
4     left = id;
5     right = (id+1) % N;
6
7     while (true) {
8         // think
9         forks[left].acquire();
10        forks[right].acquire();
11        // eat
12        forks[left].release();
13        forks[right].release();
14    }
15 }
```

## Dining Philosophers (naive attempt)

```
1 global Semaphore[] forks = [1,...,1]; // N
2
3 Philosopher(id) {
4     left = id;
5     right = (id+1) % N;
6
7     while (true) {
8         // think
9         forks[left].acquire();
10        forks[right].acquire();
11        // eat
12        forks[left].release();
13        forks[right].release();
14    }
15 }
```

**Deadlock:** If they all take the left fork, circular waiting



# Dining Philosophers (general semaphore)

```
1 global Semaphore[] forks = [1,...,1]; // N
2 global Semaphore chairs = new Semaphore(N-1);
3
4 Philosopher(id) {
5     left = id;
6     right = (id+1) % N;
7
8     while (true) {
9         // think
10        chairs.acquire();
11        forks[left].acquire();
12        forks[right].acquire();
13        // eat
14        forks[left].release();
15        forks[right].release();
16        chairs.release();
17    }
18 }
```

# Dining Philosophers (breaking the symmetry)

```
1 global Semaphore[] forks = [1,...,1]; // N
2
3 Philosopher(id) {
4     if (i == 0) {
5         left = 1;
6         right = 0;
7     } else {
8         left = id;
9         right = (id+1) % N;
10    }
11
12    while (true) {
13        // think
14        forks[left].acquire();
15        forks[right].acquire();
16        // eat
17        forks[left].release();
18        forks[right].release();
19    }
20 }
```